



NANYANG
TECHNOLOGICAL
UNIVERSITY
SINGAPORE

AI6101

Introduction to AI and AI Ethics

Reinforcement Learning

Assoc Prof Bo AN

www.ntu.edu.sg/home/boan

Email: boan@ntu.edu.sg


Office: N4-02b-55





Lesson Outline

- Intro of reinforcement learning
- Some RL algorithms
- Advanced materials:
 - Policy gradient methods
 - Exploitation vs exploration






Reinforcement Learning

What makes reinforcement learning different from other machine learning paradigms?

- There is no supervisor, only a reward signal
- Feedback is delayed, not instantaneous
- Time really matters (sequential, non i.i.d data)
- Agent's actions affect the subsequent data it receives

Examples of Reinforcement Learning



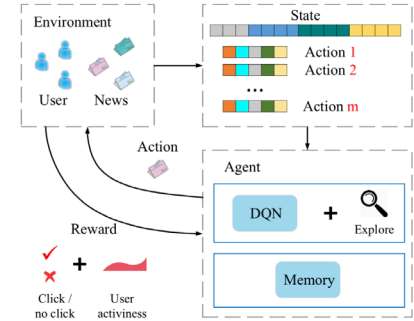
Robotic control



The Game of Go




Video games



Recommendation System



Agent and Environment



- The observation is the perception of the environment for agent
- The action will change the outside environment
- The reward is a scalar value indicates how well agent is doing at step t

The agent's job is to maximize the cumulative reward



Major Components of an RL Agent

- Policy: agent's behavior function
- Value function: how good is each state and/or action
- Model: agent's representation of the environment



Reinforcement Learning


- Motivation
 - In last lecture, we compute the value function and find the optimal policy
 - But if without the transition function $P(s'|s, a)$?
 - We can learn the value function and find the optimal policy without transition
 - From experience





RL algorithms


- Types
 - Monte Carlo
 - Q-Learning
 - DQN
 - ...





What is Monte Carlo

- Idea behind MC:
 - Just use randomness to solve a problem
- Simple definition:
 - Solve a problem by generating suitable random numbers and observing the fraction of numbers obeying some properties
- An example for calculating π (not policy in RL):
 - $S_{red} = \frac{1}{4}\pi r^2, S_{square} = r^2$
 - putting dots on the square randomly for $n = 3000$ times
 - $\pi \approx 4 \times \frac{N_{red}}{n}$, N_{red} is the number of dots in the circle






Monte Carlo in RL: Prediction

- Basic Idea: we run in the world randomly and gain experience to learn
- What experience? Many trajectories!
 - $(s_1, a_1, r_2, s_2, a_2, r_3, \dots, s_T), \dots$
- What we learn? Value function!
 - Recall that the return is the total discounted rewards:
$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^n r_{t+n} + \dots = \sum_i \gamma^i r_{t+i}$$
 - Recall that the value function is the expected return from s
$$V_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$
- How we learn?
 - Use experience to learn an empirical state value function $\tilde{V}_{\pi}(s) = \frac{1}{N} \sum_{i=1}^N G_{i,s}$



An Example

- One-dimensional grid world
 - A robot is in a 1x4 world
 - State: current cell $s \in [cell_1, cell_2, cell_3, cell_4]$
 - Action: left or right
 - Reward:
 - Move one step (-1)
 - Reach the destination cell (+10) (ignoring the one-step reward)




Start point




Destination

One-dimensional Grid World



- Trajectory or episode:

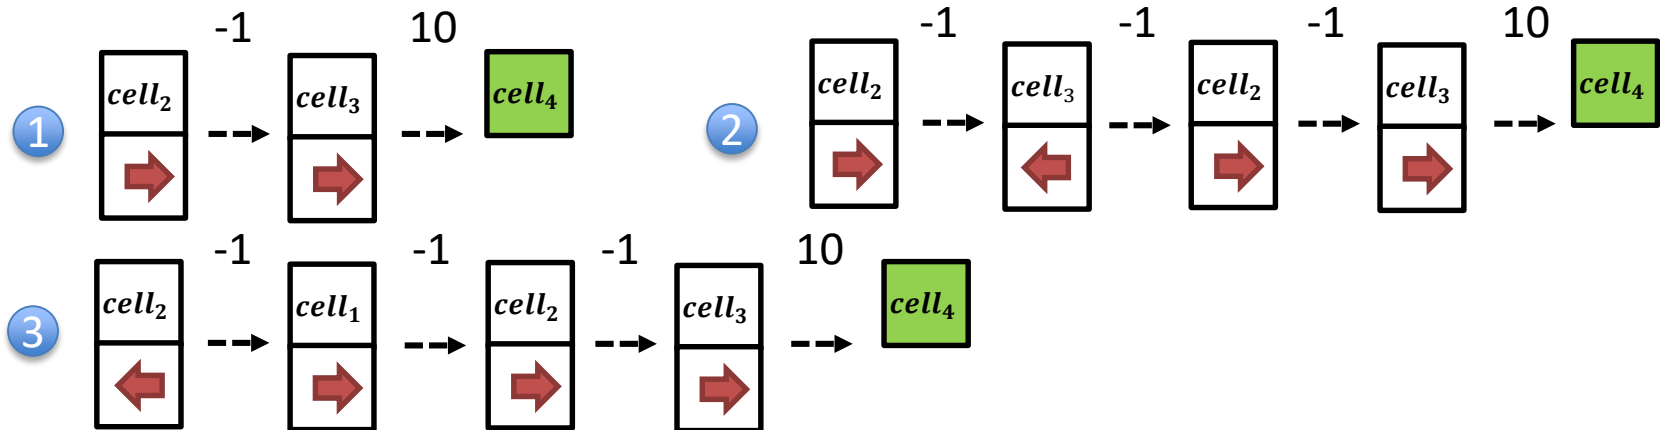
- The sequence of states from the starting state to the terminal state
- Robot starts in $cell_2$, ends in $cell_4$



Start point

Destination

- The representation of the three episodes






Compute Value Function

- Idea: Average return observed after visits to (s, a)
- First-visit MC: average returns only for **first** time (s, a) is visited in an episode
- Return in one episode (trajectory):

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^n r_{t+n} + \dots = \sum_i \gamma^i r_{t+i}$$

- We calculate the return for $cell_2$ of first episode with $\gamma = 0.9$




$$G_t = -1 \times 0.9^0 + 10 \times 0.9^1 = 8$$




Compute Value Function (cont'd)

- Similarly the return for $cell_2$ of second episode with $\gamma = 0.9$



$$G_t = -1 \times 0.9^0 - 1 \times 0.9^1 - 1 \times 0.9^2 + 10 \times 0.9^3 = 4.58$$

- Similarly the return for $cell_2$ of third episode with $\gamma = 0.9$



$$G_t = -1 \times 0.9^0 - 1 \times 0.9^1 - 1 \times 0.9^2 + 10 \times 0.9^3 = 4.58$$

- The empirical value function for $cell_2$ is $\frac{8 + 4.58 + 4.58}{3} = 5.72$



Compute Value Function (cont'd)

- Given these three episodes, we compute the value function for all non-terminal state

6.2	5.72	8.73
$cell_1$	$cell_2$	$cell_3$

- We can get more accurate value function with more episodes





First Visit Monte Carlo Policy Evaluation

- Average returns only for the first time s is visited in an episode
- Algorithm
 - Initialize:
 - $\pi \leftarrow$ policy to be evaluated
 - $V \leftarrow$ an arbitrary state-value function
 - $Returns(s) \leftarrow$ an empty list, for all state s
 - Repeat many times:
 - Generate an episode using π
 - For each state s appearing in the episode:
 - $R \leftarrow$ return following **the first occurrence** of s
 - Append R to $Returns(s)$
 - $V(s) \leftarrow average(Returns(s))$



Monte Carlo in RL: Control

- Now, we have the value function of all states given a policy
- We need to improve policy to be better
- Policy Iteration
 - Policy evaluation
 - Policy improvement
- However, we need to know how good an action is






Q-value

- Estimate how good an action is when staying in a state
- Defined as the expected return starting from s , taking the action a and thereafter following policy π

$$Q^\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$$


- Representation: A table
 - Filled with the Q-value given a state and an action





Computing Q-value


- MC for estimating Q:
 - A slight difference from estimating the value function
 - Average returns for state-action pair (s, a) is visited in an episode
- We calculate the return for $(cell_2, \text{right})$ of first episode with $\gamma = 0.9$






Compute Q-Value (cont'd)

- Similarly the return for ($cell_2$, right) of second episode with $\gamma = 0.9$



- Similarly the return for ($cell_2$, right) of third episode with $\gamma = 0.9$











- The empirical Q-value function for ($cell_2$, right) is $\frac{8 + 4.58 + 8}{3} = 6.86$



Q-Value for Control

- Filling the Q-table
 - By going through all state-action pairs, we get a complete Q-table with all the entries filled
 - A possible Q-table example

cell₁ cell₂ cell₃

	8.1	9.3	9.9
	4.5	5.6	7.5
			
			

- Selecting action

$$\pi'(s) = \operatorname{argmax}_{a \in A} Q^\pi(s, a)$$

At *cell₁*, *cell₂* and *cell₃*, we choose right



MC control algorithm

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

$\pi \leftarrow$ an arbitrary ε -soft policy

Repeat forever:

(a) Generate an episode using π

(b) For each pair s, a appearing in the episode:

$R \leftarrow$ return following the first occurrence of s, a

Append R to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

(c) For each s in the episode:

$a^* \leftarrow \arg \max_a Q(s, a)$

For all $a \in \mathcal{A}(s)$:

$$\pi(s, a) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon / |\mathcal{A}(s)| & \text{if } a = a^* \\ \varepsilon / |\mathcal{A}(s)| & \text{if } a \neq a^* \end{cases}$$

Policy evaluation

Policy improvement



Q-Learning


- Previously, we need the whole trajectory
- In Q-Learning, we only need one-step trajectory: (s, a, r, s')
- The difference is the Q-value computing

– Previously:

$$\tilde{Q}_{\pi}(s, a) = \frac{1}{N} \sum_{i=1}^N G_{i,s}$$

– Now, updating rule:

$$Q_{new}(S_t, A_t) \leftarrow Q_{old}(S_t, A_t) + \alpha (R_{t+1} + \gamma \max_a Q_{old}(S_{t+1}, a) - Q_{old}(S_t, A_t))$$





Q-Learning

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$



$S \leftarrow S'$

 until S is terminal



A Step-by-step Example


- 5-room environment as MDP
 - We'll number each room 0 through 4
 - The outside of the building can be thought of as one big room 5
 - End at room 5
 - Notice that doors at rooms 1 and 4 lead into the building from room 5 (outside)





A Step-by-step Example (cont'd)

- Goal
 - Put an agent in any room, and from that room, go outside (or room 5)
- Reward
 - The doors that lead immediately to the goal have an instant reward of 100
 - Other doors not directly connected to the target room have zero reward



$$R = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} & \text{action} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 100 \end{bmatrix} \end{matrix}$$


state



Q-Learning Step by Step

- Initialize matrix Q as a zero matrix
- $\alpha = 0.01, \gamma = 0.99$
- Loop for each episode until converge
 - Initial state: current we are in room 1 (1st outer loop)
 - Loop for each step of episode (until reach room 5)
 - ... (Next slide)

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$







Q-Learning Step by Step (cont'd)

- ... (last slide)
 - Loop for each step of episode (until room 5)
 - By random selection, we go to 5
 - We get 100 reward
 - Update Q: $Q_{new}(S_t, A_t) \leftarrow Q_{old}(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q_{old}(S_{t+1}, a) - Q_{old}(S_t, A_t))$
 - At room 5, we have 3 possible actions: go to 1, 4 or 5; We select the one with max reward
 - $Q_{new}(1,5) \leftarrow Q_{old}(1,5) + \alpha \left(100 + \gamma \max_a Q_{old}(5, a) - Q_{old}(1,5) \right) = 0 + 0.01 \times (100 + 0.99 \times 0 - 0) = 1$

$$Q = \begin{array}{c} \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} \end{array} \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



$$Q = \begin{array}{c} \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} \end{array} \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & \mathbf{1} \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$




Q-Learning Step by Step (cont'd)

- When we loop many episodes, we can get

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 80 & 0 \\ 0 & 0 & 0 & 64 & 0 & 100 \\ 0 & 0 & 0 & 64 & 0 & 0 \\ 0 & 80 & 51 & 0 & 80 & 0 \\ 64 & 0 & 0 & 64 & 0 & 100 \\ 0 & 80 & 0 & 0 & 80 & 100 \end{bmatrix} \end{matrix}$$

- According to this Q-table, we can select actions
 - E.g. We are at room 2
 - Greedily select based on maximum of Q value





An Example of Iteration Process

- A complex grid world example
- https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_td.html



Value-based Methods: SARSA

- SARSA Introduction

- Similar to Q-learning with some differences
 - **On-policy**: update the Q-table with the (s, a, r, s') samples generate by the current policy

$$Q(S, A) = Q(S, A) + \alpha \cdot [R + \gamma \cdot Q(S', A') - Q(S, A)]$$

on-policy: updating the Q with current policy

The next state and next action in transition samples

- Epsilon greedy can still be used to output actions like Q-learning

For code, refer to <https://www.geeksforgeeks.org/sarsa-reinforcement-learning/>



Value-based Methods: SARSA

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$


$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Advanced Materials



Policy Gradient Methods: Background



- How do value function work as a policy?
 - Output actions with the best Q values
- Can we directly learn a policy mapping states to actions?
- Policy gradient methods
 - Learn a *parameterized* policy that can select actions without consulting a value function
 - Use $\pi(a|s, \theta) = \Pr\{A_t = a \mid S_t = s, \theta_t = \theta\}$ with parameters $\theta \in \mathbb{R}^{d'}$ for the probability that action a is taken at time t given that the environment is in state s at time t
 - The value functions can also be parameterized as $\hat{v}(s, \mathbf{w})$ with parameters $\mathbf{w} \in \mathbb{R}^d$ (optional)
 - Update the parameters by gradient ascent given some performance measures $J(\theta)$

$$\theta_{t+1} = \theta_t + \alpha \widehat{\nabla J(\theta_t)}$$



Policy Gradient Methods: Linear Example

- A linear function approximation example of policy gradient methods

- Parameters of policy function of a linear function, **soft-max policy**

$$\theta = \{\theta_0, \theta_1, \theta_2\} \quad \text{action} = \{a_0, a_1, a_2\} \quad s = \{s_0, s_1\}$$

$$\pi(s, a_i) = \frac{e^{a_i\theta_0 + s_0\theta_1 + s_1\theta_2}}{\sum_{j=0}^{|A|} e^{a_j\theta_0 + s_0\theta_1 + s_1\theta_2}}$$

Sampling actions with these probability

- We introduce optimization rules in the following slides
- Deep Neural networks can also be used as the approximation function
 - Deep Reinforcement Learning (DRL)



Policy Gradient Methods

- Policy gradient (PG) methods model and optimize the policy directly

$$\pi_{\theta}(s, a) = \mathbb{P}[a \mid s, \theta]$$

- By maximizing performance measure w.r.t π_{θ}

$$J(\theta) = V^{\pi_{\theta}} = E[R]$$



Policy Gradient Methods

- Policy gradient (PG) methods model and optimize the policy directly
- The policy is modeled with a parameterized function respect to θ , $\pi_\theta(a|s)$

Performance measure

Transition function


$$J(\theta) = \sum_{s \in \mathcal{S}} d^\pi(s) V^\pi(s) = \sum_{s \in \mathcal{S}} d^\pi(s) \sum_{a \in \mathcal{A}} \pi_\theta(a|s) Q^\pi(s, a)$$

Value function

Gradients

$$\begin{aligned} \nabla_\theta J(\theta) &= \nabla_\theta \sum_{s \in \mathcal{S}} d^\pi(s) \sum_{a \in \mathcal{A}} Q^\pi(s, a) \pi_\theta(a|s) \\ &\propto \sum_{s \in \mathcal{S}} d^\pi(s) \sum_{a \in \mathcal{A}} Q^\pi(s, a) \nabla_\theta \pi_\theta(a|s) \end{aligned}$$

Policy Gradient Methods: REINFORCE



- REINFORCE (Monte-Carlo policy gradient) relies on an estimated return by Monte-Carlo methods using episode samples to update the policy parameter θ

$$\begin{aligned}\nabla_{\theta} J(\theta) &\propto \sum_{s \in \mathcal{S}} d^{\pi}(s) \sum_{a \in \mathcal{A}} Q^{\pi}(s, a) \nabla_{\theta} \pi_{\theta}(a|s) \\ &= \sum_{s \in \mathcal{S}} d^{\pi}(s) \sum_{a \in \mathcal{A}} \pi_{\theta}(a|s) Q^{\pi}(s, a) \frac{\nabla_{\theta} \pi_{\theta}(a|s)}{\pi_{\theta}(a|s)} \\ &= \mathbb{E}_{\pi} [Q^{\pi}(s, a) \nabla_{\theta} \ln \pi_{\theta}(a|s)] && ; \text{Because } (\ln x)' = 1/x \\ &= \mathbb{E}_{\pi} [G_t \nabla_{\theta} \ln \pi_{\theta}(A_t|S_t)] && ; \text{Because } Q^{\pi}(S_t, A_t) = \mathbb{E}_{\pi} [G_t | S_t, A_t]\end{aligned}$$

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Policy Gradient Methods: REINFORCE




1. Initialize the policy parameter θ at random.
2. Generate one trajectory on policy π_θ : $S_1, A_1, R_2, S_2, A_2, \dots, S_T$.
3. For $t=1, 2, \dots, T$:
 1. Estimate the the return G_t ;
 2. Update policy parameters: $\theta \leftarrow \theta + \alpha \gamma^t G_t \nabla_\theta \ln \pi_\theta(A_t|S_t)$

Policy Gradient Methods: Actor-Critic



- Actor-critic methods consist of two models
 - Critic updates the value function parameters w and depending on the algorithm it could be action-value $Q_w(a|s)$ or state-value $V_w(s)$
 - Actor updates the policy parameters θ for $\pi_\theta(a|s)$, in the direction suggested by the critic



1. Initialize s , θ , w at random; sample $a \sim \pi_\theta(a|s)$.
2. For $t = 1 \dots T$:
 1. Sample reward $r_t \sim R(s, a)$ and next state $s' \sim P(s'|s, a)$;
 2. Then sample the next action $a' \sim \pi_\theta(a'|s')$;
 3. Update the policy parameters: $\theta \leftarrow \theta + \alpha_\theta Q_w(s, a) \nabla_\theta \ln \pi_\theta(a|s)$;
 4. Compute the correction (TD error) for action-value at time t :
$$\delta_t = r_t + \gamma Q_w(s', a') - Q_w(s, a)$$
and use it to update the parameters of action-value function:
$$w \leftarrow w + \alpha_w \delta_t \nabla_w Q_w(s, a)$$
 5. Update $a \leftarrow a'$ and $s \leftarrow s'$.



Exploitation vs Exploration

- Online decision-making involves a fundamental choice:
 - **Exploitation** Make the best decision given current information
 - **Exploration** Gather more information
- The best long-term strategy may involve short-term sacrifices
- Gather enough information to make the best overall decisions



Examples

Restaurant Selection

- Exploitation Go to your favorite restaurant
- Exploration Try a new restaurant

Online Banner Advertisements

- Exploitation Show the most successful advert
- Exploration Show a different advert

Oil Drilling

- Exploitation Drill at the best known location
- Exploration Drill at a new location

Game Playing

- Exploitation Play the move you believe is best
- Exploration Play an experimental move



Principles

Naive Exploration

- Add noise to greedy policy (e.g. ϵ -greedy)

Optimistic Initialization

- Assume the best until proven otherwise

Optimism in the Face of Uncertainty

- Prefer actions with uncertain values

Probability Matching

- Select actions according to probability they are best

Information State Search

- Lookahead search incorporating value of information